

Enhancing Donor Prediction Accuracy for the National Veterans' Organization through Advanced Modeling Techniques

Introduction

Non-profit organizations significantly depend on donations to fund their operations and deliver services. The National Veterans' Organization (NVO), with a vast database of over 13 million potential donors, primarily utilizes direct mail for fundraising. However, with an average donation return of \$13 against a mailing cost of \$13.60 for every 20 letters sent, the organization faces a net loss. This scenario underlines the critical need for a more targeted marketing strategy that reduces costs and enhances the efficiency of fundraising efforts.

Objectives

The objective of this project is twofold:

1. Utilize various predictive modeling techniques to optimize the direct mail campaign, aiming to increase the profitability of the National Veterans' Organization.
2. Apply and showcase the modeling techniques learned in Dr. Campbell's class and beyond this semester, emphasizing practical applications in a real-world context.

Data Sources

The data was made available on Canvas by Dr. Campbell.

Details of the Data

The training data contains 3,000 records and is evenly balanced between Donors and Non Donors. There are 22 variables in the dataset. These include:

- Zip code category, homeowner status, number of children, income, sex, wealth category, homevalue, median family income, average family income, number of promotions sent to the individual, lifetime number of promotions received by the individual, sum of all donations given, dollar amount of last gift, months since last donation, number of months between previous 2 gifts, average gift from donor.

Types of analysis Performed

I performed both traditional modeling and utilized an artificial neural network to generate my predictions.

Traditional Modeling

In the traditional modeling, I started by creating dummy variables of all of the predictors. Next I removed predictor variables that were highly correlated with one another. Starting with all predictors, I would calculate the Variance Inflation Factor (Vif) of each variable, then if any variable had a Vif greater than the cut off (I used 10), I would remove the variable and recalculate the Vif without that variable. This process repeated until all Vifs were under the cut off. At the end of the process I had removed five variables


REMOVED: ['zipconvert5_Yes', 'avg_fam_inc', 'months_since_donate', 'med_fam_inc', 'avg_gift']

I fit the data to a Logistic Regression, LDA, QDA, KNN (k=1), and a Naive Bayes classifier.

The results are shown below. Our top performer on both the dev and test set was LDA. It is also noteworthy that LDA had the highest F1 score, achieving a good balance between precision and recall.

	name	tp	tn	fp	fn	acc	prec	recall	f1	test_acc
	Logistic Regression	167	151	145	137	0.530000	0.535256	0.549342	0.542208	0.533333
	LDA	163	166	149	122	0.548333	0.522436	0.571930	0.546064	0.558333
	QDA	26	268	286	20	0.490000	0.083333	0.565217	0.145251	0.525000
	KNN	148	153	164	135	0.501667	0.474359	0.522968	0.497479	0.475000
	Naïve Bayes	38	263	274	25	0.501667	0.121795	0.603175	0.202667	0.516667

To get a rough idea of what variables our model is using to generate predictions, we can make a correlation chart.



Here we can see that the larger the bar, the more correlated the variable is with the target. Red bars show negative correlations while blue bars show positive correlations. One thought I had, to achieve superior test accuracy, was to only use the most correlated variables on each side. I tested only using variables with an absolute correlation coefficient above .05. Unfortunately, This did not improve results.

Speaking of which, it is worth mentioning that I attempted many more models, and much more complex ones. I had a total of 79 entries on the competition's website. For the more complex models I tried transforming the continuous variables by log, squared, cubed, square root, inverse, boxcox, sigmoid, sine and cosine transformations. After Vif, in the complex modeling, I had a total of 84 variables. However, they did not help. Both the dev scores and test accuracy were several percentages shorter across each category.

ANN

Another approach that I took was making an artificial neural network for the classification task. This network had the following parameters:

1. An input layer with both categorical and continuous features processed via embedding and batch normalization respectively
2. Multiple hidden layers with configurations of [200, 100] neurons each incorporating ReLU activation and dropout regularization set at 0.4
3. Final output layer using a softmax function for binary classification
4. Robustly scaled input variables to boost performance
5. 100 epochs of training with an early stopping patience of 5
6. Cross validation with 20 splits

This approach achieved much better dev accuracy with an OOB accuracy of 59%. However, at test time it performed practically the same as LDA, and it had the downside of being much more complex, less interpretable, and doesn't always produce the same results.

Results

The models displayed varying levels of success:

- Accuracy Metrics: Each traditional model was evaluated based on accuracy, precision, recall, and F1-score. The LDA emerged as the most accurate model, as well as the one with the highest F1, showcasing a balance between precision and recall.
- Feature Importance: Looking at the correlation with target variable plot, we can see that the target variable is most strongly associated with number of promotions, and negatively associated with last gift. This loosely suggests that individuals who have donated recently are more likely to donate again, and people who receive a lot of promotions tend to donate more.

Discussion

The results indicated that, with all their bells and whistles, advanced modeling techniques such as artificial neural networks do not always have better predictive power than their traditional less complicated counterparts. However, none of the variables exhibited extremely high correlations with the target, suggesting potential limitations in the dataset's explanatory power.

Conclusion and Future Work

The study successfully demonstrated the potential predictive models to enhance donor prediction accuracy, thereby aiding the NVO in optimizing its fundraising strategies. For future work, further experimentation with ANN hyperparameters (such as learning rates, epochs, batch sizes, dropout rates) is recommended. Additionally, obtaining more comprehensive data, including demographic information, military service details, and any other variables that could potentially have a higher correlation with the target variable, would be useful. Having more variables more closely correlated to the target would significantly improve predictive power.

Appendices

The appendices contain two different notebooks. The first is the simple modeling I performed to achieve the best results. The second includes transformed variables (test accuracy was not included, as all results were in the mid 40s). Finally there is a python script that is my artificial neural network.

Simple Model By Hand

```
In [126]: 1 import pandas as pd
2 pd.set_option("max_colwidth", None)
3
4 import pycaret
5 import numpy as np
6 import matplotlib.pyplot as plt
7 from pycaret.classification import *
8 from sklearn.model_selection import train_test_split
9 from sklearn.metrics import accuracy_score
10
11 from functions.homebrew import *
12 import numpy as np
13 import pandas as pd
14
15 from sklearn.linear_model import LogisticRegression
16 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA, QuadraticDiscriminantAnalysis as QDA
17 from sklearn.naive_bayes import GaussianNB
18 from sklearn.neighbors import KNeighborsClassifier
19 from sklearn.preprocessing import StandardScaler
20 from sklearn.model_selection import train_test_split, cross_val_score
21 from sklearn.metrics import accuracy_score
22 from tqdm import tqdm
23 from itertools import combinations
24 import pickle
25 import os
26
27 # If you're using statsmodels or ISLP for specific tasks, keep these imports
28 import statsmodels.api as sm
29 # Assuming ISLP and homebrew are custom modules specific to your project
30 from ISLP import load_data, confusion_table
31 from ISLP.models import ModelSpec as MS, summarize, contrast
32 import statsmodels.api as sm
33 from scipy import stats
```

Helper Functions

```
In [127]: 1 def convert_confusion_matrix(df, name):
2     """
3     Converts a confusion matrix dataframe into a format with columns for model name, TP, TN, FP, FN.
4
5     Args:
6     df (pd.DataFrame): Confusion matrix dataframe with multi-index (Truth, Predicted) and columns [0, 1].
7
8     Returns:
9     pd.DataFrame: Reformatted dataframe with model evaluation metrics.
10    """
11    # Extracting the values from the confusion matrix
12    tn, fp, fn, tp = df.iloc[0, 0], df.iloc[0, 1], df.iloc[1, 0], df.iloc[1, 1]
13    acc = (tp + tn) / (tp + tn + fp + fn)
14    prec = tp / (tp + fp)
15    recall = tp / (tp + fn)
16    f1 = 2 * ((prec * recall) / (prec + recall))
17    # Creating a new dataframe with the desired format
18    metrics_df = pd.DataFrame({
19        "name": name,
20        "tp": [tp],
21        "tn": [tn],
22        "fp": [fp],
23        "fn": [fn],
24        "acc": acc,
25        "prec": prec,
26        "recall": recall,
27        "f1": f1
28    })
29
30    return metrics_df
```

```
In [128]: 1 def format_results(df):
2     df = np.where(df == 1, 'Donor', 'No Donor')
3     return df
```

LOAD DATA

```
In [140]: 1 df = pd.read_csv('./data/df.csv').drop('Unnamed: 0', axis=1)
```

```
In [141]: 1 train = df[df['type'] == 'train'].drop('type',axis =1)
2 dev = df[df['type'] == 'dev'].drop('type',axis =1)
3 test = df[df['type'] == 'test'].drop('type',axis =1)
```

VIF

```
In [142]: 1 dummies = pd.get_dummies(df, drop_first=True)
2
3 kept, removed = remove_high_vif_features(X=dummies.drop('target_No Donor', axis=1), y=dummies['target_No Donor'], vif_threshol
4 print('REMOVED:', removed)
```


REMOVED: ['avg_fam_inc', 'months_since_donate', 'zipconvert5_Yes', 'med_fam_inc', 'avg_gift']

```
In [144]: 1 kept['target'] = (df['target'] == 'Donor').astype(int)
```

```
In [145]: 1 train = kept[kept['type_train'] ==1]
2 dev = kept[(kept['type_test'] == 0) & (kept['type_train'] == 0)]
3 test = kept[kept['type_test'] ==1]
```

```
In [154]: 1 # for data in [train, dev, test]:
2 #     data.drop('type_train', inplace = True, axis = 1)
3 #     data.drop('type_test', inplace = True, axis = 1)
4 test = test.drop('target',axis = 1)
```

```
In [174]: 1 def cor_bars(df):
2     corr = df.corr()
3     # Isolating the column that represents the correlation with the target variable
4     target_corr = corr['target'].sort_values(ascending=False)
5
6     # Removing the target variable from itself to avoid a perfect correlation display
7     target_corr = target_corr.drop(labels=['target'])
8
9     # Plotting the correlations for visual representation
10    plt.figure(figsize=(10, 10))
11    sns.barplot(x=target_corr.values, y=target_corr.index, palette='coolwarm')
12    plt.title('Correlation with Target Variable (target)')
13    plt.xlabel('Correlation Coefficient')
14    plt.ylabel('Predictors')
15    plt.show()
16    cor_bars(train)
```



Logistic Regression

```
In [155]: 1 for col in kept.columns:
          2     print(col)
```

```
num_child
income
wealth
home_value
pct_lt15k
num_prom
lifetime_gifts
largest_gift
last_gift
time_lag
zipconvert2_Yes
zipconvert3_Yes
zipconvert4_Yes
homeowner_Yes
female_Yes
type_test
type_train
target
```

```
In [156]: 1 results_df = pd.DataFrame()
```

```
In [157]: 1 # Selecting features and target variable for training data
          2 X_train = train.drop(['target'], axis = 1 )
          3 y_train = train['target']
          4 X_test = dev.drop(['target'], axis = 1)
          5 y_test = dev['target']
          6
          7 # Fitting logistic regression model
          8 glm = sm.GLM(y_train, X_train, family=sm.families.Binomial())
          9 glm = glm.fit()
          10
          11 # Summarizing results
          12 # print(results.summary())
```

```
In [158]: 1 log_preds = (glm.predict(X_test) >= 0.5).astype(int)
          2 log_acc = accuracy_score(log_preds, y_test)
          3 print(log_acc)
          4
          5 d = confusion_table(log_preds,y_test)
          6 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'Logistic Regression')])
          7
          8 log_test_preds = (glm.predict(test) >= 0.5).astype(int)
          9 log_test_preds = format_results(log_test_preds)
          10
          11 save_df = pd.DataFrame(log_test_preds, columns=['values'])
          12 save_df.to_csv('./preds/log.csv', index=False)
```

0.53

LDA

```
In [159]: 1 lda = LDA(store_covariance=True)
          2 lda.fit(X_train, y_train)
          3
          4 lda_preds = lda.predict(X_test)
          5
          6 lda_acc = accuracy_score(lda_preds,y_test)
          7 print(lda_acc)
          8
          9 d = confusion_table(lda_preds,y_test)
          10 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'LDA')])
          11
          12
          13 lda_test_preds = (lda.predict(test) >= 0.5).astype(int)
          14 lda_test_preds = format_results(lda_test_preds)
          15
          16 save_df = pd.DataFrame(lda_test_preds, columns=['values'])
          17 save_df.to_csv('./preds/lda.csv', index=False)
```

0.5483333333333333

QDA.

```
In [160]: 1 qda = QDA(store_covariance=True)
2 qda.fit(X_train, y_train)
3
4 qda_preds = qda.predict(X_test)
5
6 qda_acc = accuracy_score(qda_preds,y_test)
7
8 print(qda_acc)
9
10 d = confusion_table(qda_preds,y_test)
11 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'QDA')])
12
13 qda_test_preds = (qda.predict(test) >= 0.5).astype(int)
14 qda_test_preds = format_results(qda_test_preds)
15
16 save_df = pd.DataFrame(qda_test_preds, columns=['values'])
17 save_df.to_csv('./preds/qda.csv', index=False)
```

0.49

KNN

```
In [161]: 1 df['type']
```

```
Out[161]: 0      train
1      train
2       dev
3      train
4      train
...
3115   test
3116   test
3117   test
3118   test
3119   test
Name: type, Length: 3120, dtype: object
```

```
In [162]: 1 knn1 = KNeighborsClassifier(n_neighbors=1)
2 knn1.fit(X_train, y_train)
3 knn1_pred = knn1.predict(X_test)
4 knn1_acc = accuracy_score(knn1_pred,y_test)
5
6 print(knn1_acc)
7
8 d = confusion_table(knn1_pred, y_test)
9 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'KNN')])
10
11 knn1_test_preds = (knn1.predict(test) >= 0.5).astype(int)
12 knn1_test_preds = format_results(knn1_test_preds)
13
14 save_df = pd.DataFrame(knn1_test_preds, columns=['values'])
15 save_df.to_csv('./preds/knn1.csv', index=False)
```

0.5016666666666667

NB

```
In [163]: 1 nb = GaussianNB()
2 nb.fit(X_train, y_train)
3 nb_preds = nb.predict(X_test)
4 nb_acc = accuracy_score(nb_preds,y_test)
5
6 print(nb_acc)
7 save_df = pd.DataFrame(nb_preds, columns=['values'])
8 save_df.to_csv('./preds/nb.csv', index=False)
9
10 d = confusion_table(nb_preds, y_test)
11 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'Naïve Bayes')])
12
13 nb_test_preds = (nb.predict(test) >= 0.5).astype(int)
14 nb_test_preds = format_results(nb_test_preds)
15
16 save_df = pd.DataFrame(nb_test_preds, columns=['values'])
17 save_df.to_csv('./preds/nb.csv', index=False)
```

0.5016666666666667

```
In [1]: 1 # results_df
```



```
In [168]: 1 test_acc = { # these are from running on the website
2           'log': 0.5333333,
3           'lda': 0.5583333,
4           'qda': 0.525,
5           'knn': 0.475,
6           'nb': 0.5166667,
7       }
```

```
In [169]: 1 results_df['test_acc'] = test_acc.values()
```

```
In [170]: 1 results_df
```

Out[170]:

	name	tp	tn	fp	fn	acc	prec	recall	f1	test_acc
0	Logistic Regression	167	151	145	137	0.530000	0.535256	0.549342	0.542208	0.533333
0	LDA	163	166	149	122	0.548333	0.522436	0.571930	0.546064	0.558333
0	QDA	26	268	286	20	0.490000	0.083333	0.565217	0.145251	0.525000
0	KNN	148	153	164	135	0.501667	0.474359	0.522968	0.497479	0.475000
0	Naïve Bayes	38	263	274	25	0.501667	0.121795	0.603175	0.202667	0.516667

By hand Complex

```
In [49]: 1 import pandas as pd
2 pd.set_option("max_colwidth", None)
3
4 import pycaret
5 import numpy as np
6 import matplotlib.pyplot as plt
7 from pycaret.classification import *
8 from sklearn.model_selection import train_test_split
9 from sklearn.metrics import accuracy_score
10
11 from functions.homebrew import *
12 import numpy as np
13 import pandas as pd
14
15 from sklearn.linear_model import LogisticRegression
16 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA, QuadraticDiscriminantAnalysis as QDA
17 from sklearn.naive_bayes import GaussianNB
18 from sklearn.neighbors import KNeighborsClassifier
19 from sklearn.preprocessing import StandardScaler
20 from sklearn.model_selection import train_test_split, cross_val_score
21 from sklearn.metrics import accuracy_score
22 from tqdm import tqdm
23 from itertools import combinations
24 import pickle
25 import os
26
27 # If you're using statsmodels or ISLP for specific tasks, keep these imports
28 import statsmodels.api as sm
29 # Assuming ISLP and homebrew are custom modules specific to your project
30 from ISLP import load_data, confusion_table
31 from ISLP.models import ModelSpec as MS, summarize, contrast
32 import statsmodels.api as sm
33 from scipy import stats
```

Helper Functions

```
In [50]: 1 def add_transformations(data, cont_cols):
2     for var in cont_cols:
3         data[f'log_{var}'] = np.log(data[var] + 1)
4         data[f'sq_{var}'] = data[var]**2
5         data[f'sqrt_{var}'] = np.sqrt(data[var])
6         data[f'inv_{var}'] = 1 / (data[var] + 1)
7         data[f'boxcox_{var}'], _ = stats.boxcox(data[var] + 1)
8         data[f'sigmoid_{var}'] = 1 / (1 + np.exp(-data[var]))
9         data[f'sin_{var}'] = np.sin(data[var])
10        data[f'cos_{var}'] = np.cos(data[var])
```

```
In [51]: 1 def convert_confusion_matrix(df, name):
2         """
3         Converts a confusion matrix dataframe into a format with columns for model name, TP, TN, FP, FN.
4
5         Args:
6         df (pd.DataFrame): Confusion matrix dataframe with multi-index (Truth, Predicted) and columns [0, 1].
7
8         Returns:
9         pd.DataFrame: Reformatted dataframe with model evaluation metrics.
10        """
11        # Extracting the values from the confusion matrix
12        tn, fp, fn, tp = df.iloc[0, 0], df.iloc[0, 1], df.iloc[1, 0], df.iloc[1, 1]
13        acc = (tp + tn) / (tp + tn + fp + fn)
14        prec = tp / (tp + fp)
15        recall = tp / (tp + fn)
16        f1 = 2 * ((prec * recall) / (prec + recall))
17        # Creating a new dataframe with the desired format
18        metrics_df = pd.DataFrame({
19            "name": name,
20            "tp": [tp],
21            "tn": [tn],
22            "fp": [fp],
23            "fn": [fn],
24            'acc': acc,
25            'prec': prec,
26            'recall': recall,
27            'f1': f1
28        })
29
30        return metrics_df
```

```
In [52]: 1 def format_results(df):
2         df = np.where(df == 1, 'Donor', 'No Donor')
3         return df
```

LOAD DATA

```
In [53]: 1 df = pd.read_csv('./data/df.csv').drop('Unnamed: 0', axis=1)
```

```
In [54]: 1 train = df[df['type'] == 'train'].drop('type', axis=1)
2         dev = df[df['type'] == 'dev'].drop('type', axis=1)
3         test = df[df['type'] == 'test'].drop('type', axis=1)
```

VIF

```
In [7]: 1 dummies = pd.get_dummies(df, drop_first=True)
2
3         cat_cols = [
4             'zipconvert2_Yes', 'zipconvert3_Yes', 'zipconvert4_Yes', 'boxcox_zipconvert5_Yes',
5             'homeowner_Yes', 'female_Yes', 'type_train', 'type_dev', 'type_test'
6         ]
7
8         cont_cols = [col for col in dummies.columns if col not in cat_cols + ['target']]
9         add_transformations(dummies, cont_cols)
10
11        kept, removed = remove_high_vif_features(X=dummies.drop('target_No Donor', axis=1), y=dummies['target_No Donor'], vif_threshold=10)
12        print('REMOVED:', removed)
```

REMOVED: ['num_child', 'income', 'cos_target_No Donor', 'sin_target_No Donor', 'sigmoid_target_No Donor', 'boxcox_target_No Donor', 'inv_target_No Donor', 'sqrt_target_No Donor', 'sq_target_No Donor', 'cos_zipconvert5_Yes', 'sin_zipconvert5_Yes', 'sigmoid_zipconvert5_Yes', 'boxcox_zipconvert5_Yes', 'inv_zipconvert5_Yes', 'sqrt_zipconvert5_Yes', 'sq_zipconvert5_Yes', 'log_zipconvert5_Yes', 'log_income', 'sq_income', 'log_num_child', 'sq_num_child', 'sigmoid_avg_fam_inc', 'inv_num_child', 'sqrt_months_since_donate', 'sqrt_med_fam_inc', 'sq_wealth', 'boxcox_avg_fam_inc', 'boxcox_num_prom', 'log_last_gift', 'months_since_donate', 'log_avg_gift', 'inv_home_value', 'inv_avg_fam_inc', 'boxcox_time_lag', 'log_wealth', 'inv_med_fam_inc', 'log_largest_gift', 'boxcox_home_value', 'sqrt_num_prom', 'log_lifetime_gifts', 'sqrt_income', 'boxcox_pct_lt15k', 'sqrt_wealth', 'sq_months_since_donate', 'sqrt_avg_fam_inc', 'sqrt_avg_gift', 'boxcox_med_fam_inc', 'sqrt_time_lag', 'boxcox_largest_gift', 'sqrt_pct_lt15k', 'inv_income', 'sqrt_home_value', 'sigmoid_num_child', 'log_num_prom', 'sqrt_last_gift', 'log_months_since_donate', 'boxcox_avg_gift', 'sqrt_lifetime_gifts', 'avg_fam_inc', 'wealth', 'largest_gift', 'med_fam_inc', 'boxcox_last_gift', 'inv_pct_lt15k', 'zipconvert5_Yes', 'num_prom', 'pct_lt15k', 'home_value', 'log_avg_fam_inc', 'inv_wealth', 'log_time_lag', 'boxcox_lifetime_gifts', 'sqrt_num_child', 'avg_gift', 'inv_time_lag', 'log_med_fam_inc', 'boxcox_income', 'inv_largest_gift', 'inv_last_gift', 'sq_avg_fam_inc', 'inv_num_prom', 'last_gift']

```
In [55]: 1 final_vars = list(kept.corr().drop('target')[np.abs(kept.corr()['target']).drop('target') > .05].index)
```

```
In [56]: 1 regress = kept[final_vars]
```

```
In [57]: 1 regress['target'] = (df['target'] == 'Donor').astype(int)
```

```
In [72]: 1 kept = kept.drop('log_target_No Donor', axis =1)
```

```
In [59]: 1 train = kept[regress['type_train'] ==1]
2 dev = kept[(regress['type_test'] == 0) & (regress['type_train'] == 0)]
3 test = kept[regress['type_test'] ==1]
```

```
In [60]: 1 for data in [train, dev, test]:
2     data.drop('type_train', inplace = True, axis = 1)
3     data.drop('type_test', inplace = True, axis = 1)
4 test = test.drop('target', axis =1)
```

Logistic Regression

```
In [61]: 1 results_df = pd.DataFrame()
```

```
In [62]: 1 # Selecting features and target variable for training data
2 X_train = train.drop(['target'], axis =1 )
3 y_train = train['target']
4 X_test = dev.drop(['target'], axis = 1)
5 y_test = dev['target']
6
7 # Fitting logistic regression model
8 glm = sm.GLM(y_train, X_train, family=sm.families.Binomial())
9 glm = glm.fit()
10
11 # Summarizing results
12 # print(results.summary())
```

```
In [63]: 1 log_preds = (glm.predict(X_test) >= 0.5).astype(int)
2 log_acc = accuracy_score(log_preds, y_test)
3 print(log_acc)
4
5 d = confusion_table(log_preds,y_test)
6 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'Logistic Regression')])
7
8 log_test_preds = (glm.predict(test) >= 0.5).astype(int)
9 log_test_preds = format_results(log_test_preds)
10
11 save_df = pd.DataFrame(log_test_preds, columns=['values'])
12 save_df.to_csv('./preds/log.csv', index=False)
```

0.5116666666666667

LDA

```
In [64]: 1 lda = LDA(store_covariance=True)
2 lda.fit(X_train, y_train)
3
4 lda_preds = lda.predict(X_test)
5
6 lda_acc = accuracy_score(lda_preds,y_test)
7 print(lda_acc)
8
9 d = confusion_table(lda_preds,y_test)
10 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'LDA')])
11
12
13 lda_test_preds = (lda.predict(test) >= 0.5).astype(int)
14 lda_test_preds = format_results(lda_test_preds)
15
16 save_df = pd.DataFrame(lda_test_preds, columns=['values'])
17 save_df.to_csv('./preds/lda.csv', index=False)
```

0.5066666666666667

QDA.

```
In [65]: 1 qda = QDA(store_covariance=True)
2 qda.fit(X_train, y_train)
3
4 qda_preds = qda.predict(X_test)
5
6 qda_acc = accuracy_score(qda_preds,y_test)
7
8 print(qda_acc)
9
10 d = confusion_table(qda_preds,y_test)
11 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'QDA')])
12
13 qda_test_preds = (qda.predict(test) >= 0.5).astype(int)
14 qda_test_preds = format_results(qda_test_preds)
15
16 save_df = pd.DataFrame(qda_test_preds, columns=['values'])
17 save_df.to_csv('./preds/qda.csv', index=False)

0.49166666666666664
```

KNN

```
In [66]: 1 knn1 = KNeighborsClassifier(n_neighbors=1)
2 knn1.fit(X_train, y_train)
3 knn1_pred = knn1.predict(X_test)
4 knn1_acc = accuracy_score(knn1_pred,y_test)
5
6 print(knn1_acc)
7
8 d = confusion_table(knn1_pred, y_test)
9 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'KNN')])
10
11 knn1_test_preds = (knn1.predict(test) >= 0.5).astype(int)
12 knn1_test_preds = format_results(knn1_test_preds)
13
14 save_df = pd.DataFrame(knn1_test_preds, columns=['values'])
15 save_df.to_csv('./preds/knn1.csv', index=False)

0.48166666666666667
```

NB

```
In [67]: 1 nb = GaussianNB()
2 nb.fit(X_train, y_train)
3 nb_preds = nb.predict(X_test)
4 nb_acc = accuracy_score(nb_preds,y_test)
5
6 print(nb_acc)
7 save_df = pd.DataFrame(nb_preds, columns=['values'])
8 save_df.to_csv('./preds/nb.csv', index=False)
9
10 d = confusion_table(nb_preds, y_test)
11 results_df = pd.concat([results_df,convert_confusion_matrix(d, 'Naïve Bayes')])
12
13 nb_test_preds = (nb.predict(test) >= 0.5).astype(int)
14 nb_test_preds = format_results(nb_test_preds)
15
16 save_df = pd.DataFrame(nb_test_preds, columns=['values'])
17 save_df.to_csv('./preds/nb.csv', index=False)

0.495
```

```
In [68]: 1 results_df
```

Out[68]:

	name	tp	tn	fp	fn	acc	prec	recall	f1
0	Logistic Regression	159	148	153	140	0.511667	0.509615	0.531773	0.520458
0	LDA	156	148	156	140	0.506667	0.500000	0.527027	0.513158
0	QDA	26	269	286	19	0.491667	0.083333	0.577778	0.145658
0	KNN	140	149	172	139	0.481667	0.448718	0.501792	0.473773
0	Naïve Bayes	18	279	294	9	0.495000	0.057692	0.666667	0.106195

```
In [69]: 1 # test_acc = {
2 #       'Log': 0.5333333,
3 #       'Lda': 0.5583333,
4 #       'qda': 0.525,
5 #       'knn': 0.475,
6 #       'nb': 0.5166667,
7 # }
```

```
In [70]: 1 # results_df['test_acc'] = test_acc.values()
```

```
In [71]: 1 results_df
```

Out[71]:

	name	tp	tn	fp	fn	acc	prec	recall	f1	test_acc
0	Logistic Regression	159	148	153	140	0.511667	0.509615	0.531773	0.520458	0.533333
0	LDA	156	148	156	140	0.506667	0.500000	0.527027	0.513158	0.558333
0	QDA	26	269	286	19	0.491667	0.083333	0.577778	0.145658	0.525000
0	KNN	140	149	172	139	0.481667	0.448718	0.501792	0.473773	0.475000
0	Naïve Bayes	18	279	294	9	0.495000	0.057692	0.666667	0.106195	0.516667

```

1 # =====
2 # SET-UP AND IMPORTS
3 # =====
4 import os
5 import torch
6 import torch.nn as nn
7 import numpy as np
8 import pandas as pd
9 import matplotlib.pyplot as plt
10 from torch.utils.data import DataLoader, TensorDataset, ConcatDataset
11 from sklearn.model_selection import KFold
12 from sklearn.preprocessing import StandardScaler
13
14 # Set directory and device setup
15 os.chdir('/home/dan/FUNDRAISING')
16 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
17 print(f"Using device: {device}")
18
19 # =====
20 # PREPARE DATA
21 # =====
22 df = pd.read_csv('./data/df.csv', index_col='Unnamed: 0')
23 df['zipconvert'] = df[['zipconvert2', 'zipconvert3', 'zipconvert4', 'zipconvert5']].apply(
24     lambda row: 'zc' + str(row.idxmax()[-1]) if pd.notna(row.idxmax()) else 'unknown', axis=1
25 )
26 test = df[df['type']=='test']
27
28 df = df[df['type']!='test']
29
30 df.drop(['zipconvert2', 'zipconvert3', 'zipconvert4', 'zipconvert5', 'type'], axis=1, inplace=True)
31 test.drop(['zipconvert2', 'zipconvert3', 'zipconvert4', 'zipconvert5', 'type'], axis=1, inplace=True)
32
33 cat_cols = ['homeowner', 'female', 'zipconvert', 'wealth', 'income', 'num_child']
34 cont_cols = [col for col in df.columns if col not in cat_cols + ['target']]
35 emb_sizes = [(df[col].astype('category').cat.codes.max() + 1, min(50, (df[col].nunique() + 1) // 2)) for col in cat_cols]
36
37 # Function to process datasets
38 def process_data(data):
39     cats = np.stack([data[col].astype('category').cat.codes.values for col in cat_cols], axis=1)
40     conts = np.stack([data[col].values for col in cont_cols], axis=1)
41     scaler = StandardScaler()
42
43     conts = scaler.fit_transform(conts)
44     y = data['target'].map({'Donor': 1, 'No Donor': 0}).values
45     return torch.tensor(cats, dtype=torch.int64), torch.tensor(conts, dtype=torch.float), torch.tensor(y, dtype=torch.long)
46
47 cats, conts, targets = process_data(df)
48 test_cats, test_conts, test_targets = process_data(test)
49
50 dataset = TensorDataset(cats, conts, targets)
51
52 test_dataset = TensorDataset(test_cats, test_conts, test_targets)
53 test_loader = DataLoader(test_dataset, batch_size=2048, shuffle=False)
54
55 # =====
56 # DEFINE A TABULAR MODEL
57 # =====
58
59 class TabularModel(nn.Module):
60     def __init__(self, emb_sizes, n_cont, out_sz, layers, p=0.5):
61         super().__init__()
62         self.embs = nn.ModuleList([nn.Embedding(ni, nf) for ni, nf in emb_sizes])
63         self.emb_drop = nn.Dropout(p)
64         self.bn_cont = nn.BatchNorm1d(n_cont)
65
66         # Calculate the total embedding output size
67         total_emb_size = sum(nf for _, nf in emb_sizes)
68         n_in = total_emb_size + n_cont # Total input size for the first linear layer
69
70         layerlist = nn.ModuleList()
71         for output_size in layers:
72             layerlist.append(nn.Linear(n_in, output_size))
73             layerlist.append(nn.ReLU())
74             layerlist.append(nn.BatchNorm1d(output_size))
75             layerlist.append(nn.Dropout(p))
76             n_in = output_size # Update n_in to the output size of the current layer
77
78         layerlist.append(nn.Linear(layers[-1], out_sz)) # Final output layer
79         self.layers = nn.Sequential(*layerlist)
80
81     def forward(self, x_cat, x_cont):
82         embeddings = [e(x_cat[:, i]) for i, e in enumerate(self.embs)]
83         x = torch.cat(embeddings, 1)
84         x = self.emb_drop(x)

```

```

85         x_cont = self.bn_cont(x_cont)
86         x = torch.cat([x, x_cont], 1)
87         return self.layers(x)
88
89 # Initialize model
90 model = TabularModel(emb_sizes, len(cont_cols), 2, [200, 100], p=0.4)
91 model = model.to(device)
92
93
94 #endregion
95 #region # EARLY STOPPING
96 # =====
97 # EARLY STOPPING
98 # =====
99 class EarlyStopping:
100     def __init__(self, patience=5, verbose=False, delta=0):
101         self.patience = patience
102         self.verbose = verbose
103         self.delta = delta
104         self.best_score = None
105         self.early_stop = False
106         self.counter = 0
107
108     def __call__(self, val_loss, model):
109         score = -val_loss
110
111         if self.best_score is None:
112             self.best_score = score
113         elif score < self.best_score + self.delta:
114             self.counter += 1
115             if self.verbose:
116                 print(f'EarlyStopping counter: {self.counter} out of {self.patience}')
117             if self.counter >= self.patience:
118                 self.early_stop = True
119         else:
120             self.best_score = score
121             self.counter = 0
122
123
124 # =====
125 # CROSS-VALIDATION SETUP
126 # =====
127 def calculate_accuracy(y_pred, y_true):
128     y_pred_classes = torch.argmax(y_pred, dim=1)
129     correct = (y_pred_classes == y_true).float() # convert into float for division
130     acc = correct.sum() / len(correct)
131     return acc
132
133 kf = KFold(n_splits=20, shuffle=True, random_state=42)
134 early_stopping = EarlyStopping(patience=5, verbose=True)
135 results = []
136
137 for fold, (train_idx, val_idx) in enumerate(kf.split(dataset)):
138     train_subsampler = torch.utils.data.SubsetRandomSampler(train_idx)
139     val_subsampler = torch.utils.data.SubsetRandomSampler(val_idx)
140     train_loader = DataLoader(dataset, batch_size=32, sampler=train_subsampler)
141     val_loader = DataLoader(dataset, batch_size=32, sampler=val_subsampler)
142
143     model = TabularModel(emb_sizes, len(cont_cols), 2, [100, 200], p=0.4).to(device)
144     criterion = nn.CrossEntropyLoss()
145     optimizer = torch.optim.Adam(model.parameters(), lr=1e-5)
146
147     for epoch in range(50): # Adjust as needed
148         model.train()
149         total_loss, total_acc = 0, 0
150         for cats, conts, y in train_loader:
151             cats, conts, y = cats.to(device), conts.to(device), y.to(device)
152             optimizer.zero_grad()
153             outputs = model(cats, conts)
154             loss = criterion(outputs, y)
155             acc = calculate_accuracy(outputs, y)
156             total_loss += loss.item()
157             total_acc += acc.item()
158             loss.backward()
159             optimizer.step()
160         avg_train_loss = total_loss / len(train_loader)
161         avg_train_acc = total_acc / len(train_loader)
162
163         model.eval()
164         val_loss, val_acc = 0, 0
165         with torch.no_grad():
166             for cats, conts, y in val_loader:
167                 cats, conts, y = cats.to(device), conts.to(device), y.to(device)
168                 outputs = model(cats, conts)
169                 loss = criterion(outputs, y)
170                 acc = calculate_accuracy(outputs, y)
171                 val_loss += loss.item()

```



```

172         val_acc += acc.item()
173     avg_val_loss = val_loss / len(val_loader)
174     avg_val_acc = val_acc / len(val_loader)
175
176     print(f'Fold {fold+1}, Epoch {epoch+1}, Train Loss: {avg_train_loss:.4f}, Train Acc: {avg_train_acc:.4f}, Val Loss: {avg_val_loss:.4f}, Val
Acc: {avg_val_acc:.4f}')
177
178     # Call early stopping
179     early_stopping(avg_val_loss, model)
180     if early_stopping.early_stop:
181         print("Early stopping")
182         break
183
184     results.append((avg_val_loss, avg_val_acc))
185
186 average_val_loss = sum(x[0] for x in results) / len(results)
187 average_val_acc = sum(x[1] for x in results) / len(results)
188 print(f'Average Validation Loss: {average_val_loss:.4f}, Average Validation Accuracy: {average_val_acc:.4f}')
189
190 preds = []
191 model.eval()
192
193 with torch.no_grad():
194     for cats, conts, y in test_loader:
195         cats, conts, y = cats.to(device), conts.to(device), y.to(device)
196         output = model(cats, conts)
197         predicted = output.argmax(dim=1) # Ensure you use dim=1
198         # print(predicted)
199         preds.append(predicted.cpu()) # Move predictions to CPU
200 # print(type(preds[0]))
201 # Concatenate all batch predictions into a single tensor
202 # preds = torch.cat(preds)
203 preds = preds[0].tolist()
204 # print(preds)
205 final_preds = []
206 donor=1
207 no_donor=1
208 for i in preds:
209     if i == 1:
210         final_preds.append('Donor')
211         donor +=1
212     else:
213         final_preds.append('No Donor')
214         no_donor+=1
215 print(f'% Donor = {donor / (donor+no_donor)}')
216
217
218 save_df = pd.DataFrame(final_preds, columns=['values'])
219 save_df.to_csv('./preds/preds.csv', index=False)

```